




Intro to classification - Logistic regression - 3

One should look for what is and not what he thinks should be. (Albert Einstein)


Warm up: overfitting

- Based on the graphs below, which model is **overfitted**? Which is **underfitted**? Which appears to be a **good fit**?
- **Share your responses** in the chat or aloud


A.



B.



C.



Source

Module completion checklist

Objectives	Complete
Analyze the model to determine if / when overfitting occurs	
Demonstrate how to tune the model using grid search cross-validation	

Accuracy on train vs. accuracy on test

- Our accuracy score for test data was:

```
0.9458577951728636
```

- Take a look at the accuracy score for the training data

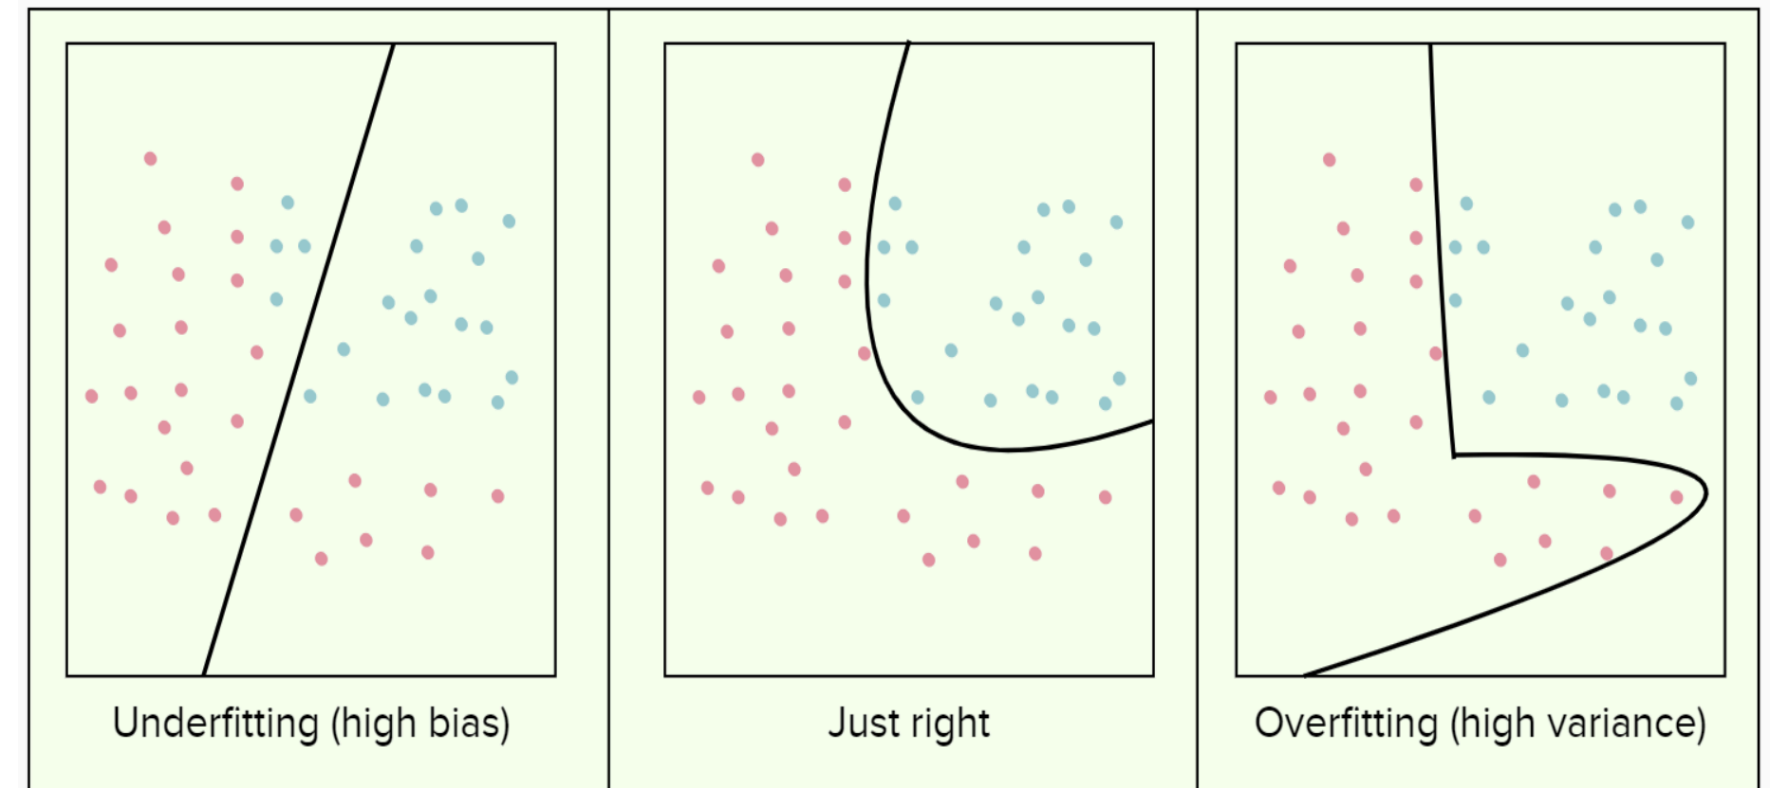
```
# Compute trained model accuracy score.  
trained_accuracy_score = logistic_regression_model.score(X_train_scaled, y_train)  
print("Accuracy on train data: ", trained_accuracy_score)
```

```
Accuracy on train data: 0.9535923958624546
```

- Did our model underperform?
- Is there a big difference in `train` and `test` accuracy?
- If there is a difference, the problem usually lies in **overfitting**


When overfitting occurs

- An overfitted model consists of high variance and usually shows a drastically higher accuracy in the training data because it doesn't generalize well to new data
- Creating a model that fits training data too well will lead to poor generalization and poor performance on new data



When overfitting occurs (cont'd)

- A model might treat **noise** as actual artifacts of the data, so when it encounters new data with new noise, it underperforms
- It might use **too many predictors** that only contribute tiny portions to variation in our data
- The train set might not be an **accurate representation** of the data, but a partial and inaccurate sample that doesn't translate



How to overcome overfitting

- Use so-called **soft-margin classifiers** to:
 - Utilize penalization constants and methods to make the model less prone to noise
 - Tune them to use the optimal parameters for best model performance
- Use **feature selection** and/ or **feature extraction** methods to:
 - Capture a few main features responsible for the most variation in the data
 - Discard the features that don't account for variation in the data
- **Gather more data**

Tuning logistic regression model

- Recall the two parameters that we mentioned before:
 - `penalty`: a regularization technique used to tune the model (either `l1`, a.k.a. *Lasso*, or `l2`, a.k.a. *Ridge*; default is `l2`)
 - `C`: a regularization constant used to amplify the effect of the regularization method (a value between $[0, \infty]$; default is `1`)
- These two parameters control a so-called **regularization term** that adds a penalty as the model complexity increases with added variables
- These two parameters play a key role in mitigating overfitting and feature pruning

Regularization techniques in logistic regression

- As you may know, any ML algorithm optimizes some cost function $f(x)$
- In logistic regression, ℓ_1 (Lasso) adds a term to that function like so:

$$f(x) + C \sum_{j=1}^n |b_j|$$

- While ℓ_2 (Ridge) adds a term like so:

$$f(x) + C \sum_{j=1}^n b_j^2$$

- You can see that Lasso uses the absolute value b_j , while Ridge uses a squared b_j
- That term, when added to the original cost function, **dampens** the margins of our classifier, making it more **forgiving** of the misclassification of some points that might be noise

Lasso vs. Ridge

Lasso (11)

$$C \sum_{j=1}^n |b_j|$$

- Stands for **L**east **A**bsolute **S**hrinkage and **S**election **O**perator
- It adds “absolute value of magnitude” of the coefficient as a penalty term to the loss function
- **Shrinks** (as the name suggests) the less important features’ coefficients to zero, which leads to **removal** of some features

Ridge (12)

$$C \sum_{j=1}^n b_j^2$$

- Adds “squared magnitude” of coefficient as penalty term to the loss function
- **Dampens** the less important features’ coefficients making them less significant, which leads to **weighting** of the features according to their importance

What is the role of C ?

There are **4 scenarios** that might happen with a classifier with respect to C :

- Scenario 1: $C = 0$
 - The classifier becomes an **OLS** problem (i.e., Ordinary Least Squares, or just a strict regression without any penalization)
 - Since $0 \times \textit{anything} = 0$, we are just left with optimizing $f(x)$, which is a definite **overfitting** problem
- Scenario 2: $C = \textit{small}$
 - We still run into an **overfitting** problem
 - Since C will not “magnify” the effect of the penalty term enough

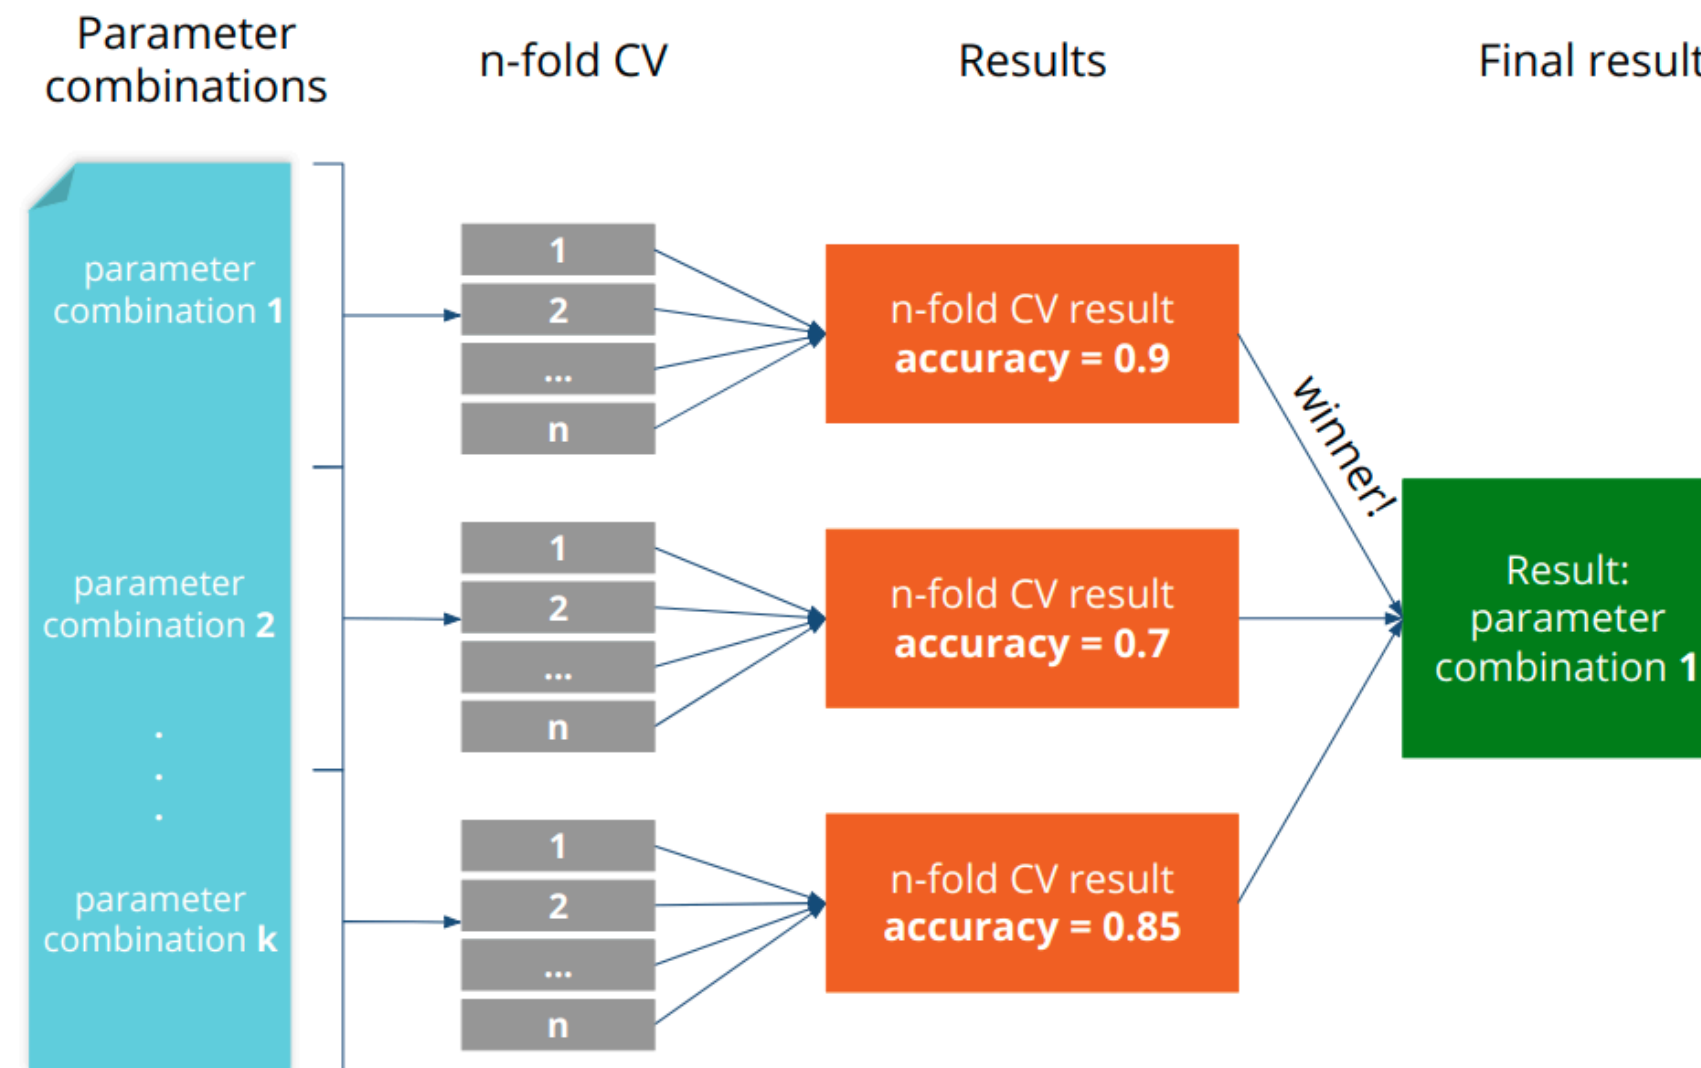
What is the role of C? (cont'd)

- Scenario 3: $C = \text{large}$
 - We run into an **underfitting** problem, where we've weighted and dampened the coefficients too much and we made the model too general
- Scenario 4: $C = \text{optimal}$
 - We have a **good, robust, and generalizable model** that works well with new data
 - The model ignores most of the noise while preserving the main pattern in data
- We can pick the right combination of parameters using a technique called **grid search cross-validation**

Module completion checklist

Objectives	Complete
Analyze the model to determine if / when overfitting occurs	✓
Demonstrate how to tune the model using grid search cross-validation	

What does grid search cross-validation do?



scikit-learn: model_selection.GridSearchCV

- `estimator` is the name of sklearn algorithm to optimize
- `param_grid` is a dictionary or list of parameters to optimize
- `cv` is an `int` of `n` for `n-fold` cross-validation
- `verbose` is an `int` of how much verbosity in messages you want to see as the function runs

sklearn.model_selection.GridSearchCV

```
class sklearn.model_selection. GridSearchCV (estimator, param_grid, scoring=None, fit_params=None,
n_jobs=1, iid=True, refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs', error_score='raise',
return_train_score='warn')
```

[\[source\]](#)

Exhaustive search over specified parameter values for an estimator.

Important members are `fit`, `predict`.

`GridSearchCV` implements a “fit” and a “score” method. It also implements “predict”, “predict_proba”, “decision_function”, “transform” and “inverse_transform” if they are implemented in the estimator used.

The parameters of the estimator used to apply these methods are optimized by cross-validated grid-search over a parameter grid.

[Click here for full documentation](#)

Prepare parameters for optimization

```
# Create regularization penalty space.  
penalty = ['l1', 'l2']
```

```
# Create regularization constant space.  
C = np.logspace(0, 10, 10)  
print("Regularization constant: ", C)
```

```
Regularization constant:  [1.00000000e+00  1.29154967e+01  1.66810054e+02  2.15443469e+03  
 2.78255940e+04  3.59381366e+05  4.64158883e+06  5.99484250e+07  
 7.74263683e+08  1.00000000e+10]
```

```
# Create hyperparameter options dictionary.  
hyperparameters = dict(C = C, penalty = penalty)  
print(hyperparameters)
```

```
{'C': array([1.00000000e+00, 1.29154967e+01, 1.66810054e+02, 2.15443469e+03,  
 2.78255940e+04, 3.59381366e+05, 4.64158883e+06, 5.99484250e+07,  
 7.74263683e+08, 1.00000000e+10]), 'penalty': ['l1', 'l2']}
```


Set up cross-validation logistic function

```
# Grid search 10-fold cross-validation with above parameters.
clf = GridSearchCV(linear_model.LogisticRegression(solver='liblinear'), #<- function to optimize
                  hyperparameters,                                     #<- grid search parameters
                  cv = 10,                                           #<- 10-fold cv
                  verbose = 0)                                       #<- no messages to show
```

```
# Fit CV grid search.
best_model = clf.fit(X_train_scaled, y_train)
best_model
```

```
GridSearchCV(cv=10, estimator=LogisticRegression(solver='liblinear'),
             param_grid={'C': array([1.00000000e+00, 1.29154967e+01, 1.66810054e+02,
2.15443469e+03,
2.78255940e+04, 3.59381366e+05, 4.64158883e+06, 5.99484250e+07,
7.74263683e+08, 1.00000000e+10]),
             'penalty': ['l1', 'l2']})
```

Check best parameters found by CV

```
# Get best penalty and constant parameters.  
penalty = best_model.best_estimator_.get_params() ['penalty']  
constant = best_model.best_estimator_.get_params() ['C']  
print('Best penalty: ', penalty)
```

```
Best penalty: 12
```

```
print('Best C: ', constant)
```

```
Best C: 1.0
```

- It seems like our grid search CV has found that 11 (i.e. *Lasso* regularization method) works better than the default 12 (i.e. *Ridge*)
- It also shows that the default C, which is 1, creates a big enough soft margin for our classifier

Predict using the best model parameters

- Now let's use the tuned model to predict on our test data

```
# Predict on test data using best model.  
best_predicted_values = best_model.predict(X_test_scaled)  
print(best_predicted_values)
```

```
[False False False ... False False False]
```

```
# Compute best model accuracy score.  
best_accuracy_score = metrics.accuracy_score(y_test, best_predicted_values)  
print("Accuracy on test data (best model): ", best_accuracy_score)
```

```
Accuracy on test data (best model):  0.9458577951728636
```

Accuracy on train vs. accuracy on test

- Take a look at the accuracy score for the training data

```
# Compute trained model accuracy score.  
trained_accuracy_score = best_model.score(X_train_scaled, y_train)  
print("Accuracy on train data: " , trained_accuracy_score)
```

```
Accuracy on train data: 0.9535923958624546
```

- What do you notice by comparing the train and test accuracy?

Assessing the tuned model

- Now we can start to evaluate this model and compare how it works as compared to the previous version

```
# Compute confusion matrix for best model.
best_confusion_matrix = metrics.confusion_matrix(y_test, best_predicted_values)
print(best_confusion_matrix)
```

```
[[1450    0]
 [  83    0]]
```

```
# Create a list of target names to interpret class assignments.
target_names = ['Low value', 'High value']
```

```
# Compute classification report for best model.
best_class_report = metrics.classification_report(y_test, best_predicted_values,
                                                  target_names = target_names)
```

```
/opt/conda/envs/sdaia-python-classification/lib/python3.7/site-
packages/sklearn/metrics/_classification.py:1245: UndefinedMetricWarning: Precision and F-
score are ill-defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/opt/conda/envs/sdaia-python-classification/lib/python3.7/site-
packages/sklearn/metrics/_classification.py:1245: UndefinedMetricWarning: Precision and F-
score are ill-defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
  warn_prf(average, modifier, msg_start, len(result))
```

Assessing the tuned model (cont'd)

```
print(best_class_report)
```

	precision	recall	f1-score	support
Low value	0.95	1.00	0.97	1450
High value	0.00	0.00	0.00	83
accuracy			0.95	1533
macro avg	0.47	0.50	0.49	1533
weighted avg	0.89	0.95	0.92	1533

Save accuracy score

- Let's save our final model

```
model_final = {'metrics' : "accuracy",  
               'values' : round(best_accuracy_score, 4),  
               'model': 'logistic_tuned' }  
  
print(model_final)
```

```
{'metrics': 'accuracy', 'values': 0.9459, 'model': 'logistic_tuned'}
```

Get metrics for ROC curve

```
# Get probabilities instead of predicted values.  
best_test_probabilities = best_model.predict_proba(X_test_scaled)  
print(best_test_probabilities[0:5, ])
```

```
[[0.97142844 0.02857156]  
 [0.83103002 0.16896998]  
 [0.98592979 0.01407021]  
 [0.85404287 0.14595713]  
 [0.96287726 0.03712274]]
```

```
# Get probabilities of test predictions only.  
best_test_predictions = best_test_probabilities[:, 1]  
print(best_test_predictions[0:5])
```

```
[0.02857156 0.16896998 0.01407021 0.14595713 0.03712274]
```


Get metrics for ROC curve (cont'd)

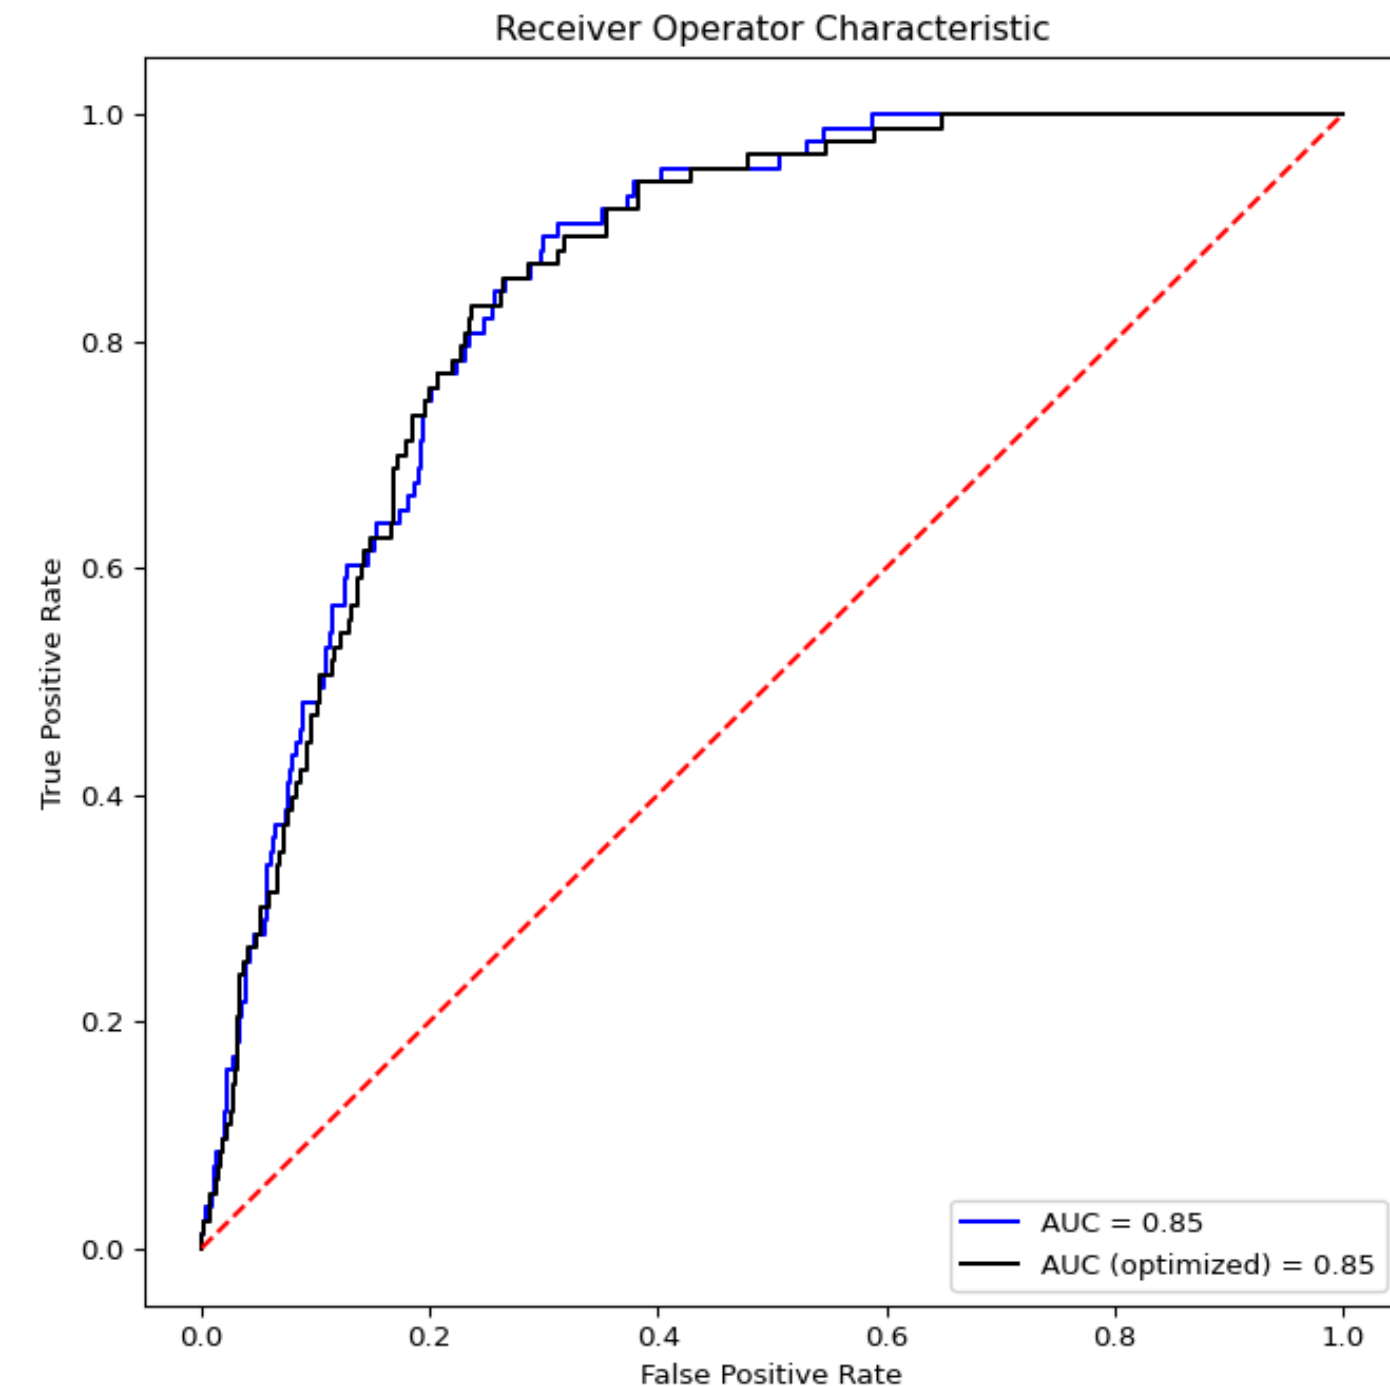
```
# Get ROC curve metrics.  
best_fpr, best_tpr, best_threshold = metrics.roc_curve(y_test, best_test_predictions)  
best_auc = metrics.auc(best_fpr, best_tpr)  
print(best_auc)
```

```
0.8516078105525551
```

Plot ROC curve for both models

```
# Make an ROC curve plot.
plt.title('Receiver Operator Characteristic')
plt.plot(fpr, tpr, 'blue',
         label = 'AUC = %0.2f'%auc)
plt.plot(best_fpr, best_tpr, 'black',
         label = 'AUC (best) = %0.2f'%best_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```


- Does it look as though our model has improved?



Knowledge check



Exercise



You are now ready to try Tasks 9-13 in the Exercise for this topic

Module completion checklist

Objectives	Complete
Analyze the model to determine if / when overfitting occurs	✓
Demonstrate how to tune the model using grid search cross-validation	✓

Logistic regression: Topic summary

In this part of the course, we have covered:

- Logistic regression use cases and theory behind it
- Data transformation necessary for logistic regression
- Implementation of logistic regression on a dataset
- Model performance evaluation and tuning

Congratulations on completing this module!

